

Enterprise agents can act. Can they be trusted?

Prompt Injection Firewall · Problem Statement #2

TRUSTBOUNDARY

AI

THE AGENTIC SECURITY CONTROL PLANE

Detect hostile instructions · Contain unauthorized tool actions · Measure true task utility

Participant: Dhinesh Babu Venkatesan

Prompt injection becomes business risk when AI can act.

An email, retrieved page or tool response may contain a forged high-priority instruction.

01 / ATTACK

Untrusted input

Merchant email hides a role override or encoded instruction.

02 / ESCALATION

Tool manipulation

Agent is pushed toward refund, secret or outbound-message tool.

03 / CONSEQUENCE

Business harm

Unauthorized financial action, leaked secrets, or interrupted workflow.

Product principle: External content supplies data, never authority.

Not a chatbot. A security layer around agent actions.

Three bounded components plus independent authorization and EGO measurement.

INSPECT

Normalize sources, inspect risky spans, assign provenance & attack categories.

ENFORCE

Deterministic role + source-trust gate; privileged tools need separate approval.

EVALUATE

Audit tool outcomes, triage incidents and replay reproducible security tests.

DEFENSE-IN-DEPTH

9 categories

Optional hosted LLM + locally trained TF-IDF risk model + rule layer; authorization stays outside the model.

Live attack laboratory and case comparison.

Independent tool gate protects user task even when classification misses.

PLATFORM / ATTACK LAB

OFFLINE HYBRID ML

CONTROL PLANE ACTIVE

01 / ADVERSARIAL SIMULATION

Attack Lab.

Inspect untrusted content, execute the protected workflow, and compare real sandbox outcomes against a deliberately insecure reference.

DETECTION TAXONOMY

9 / 9 implemented

Heuristic patterns; coverage varies

ENFORCEMENT LAYER

Zero Trust

Deterministic tool authorization

WORKFLOW CONTROLS

3 agents

Inspection · Execution · Incident

LIVE SANDBOX

TRUSTBOUNDARY

AGENTIC SECURITY

WORKSPACE

Attack Lab

Evaluation Cockpit

Audit & Traces

01 / Input source

LOW-TRUST CONTENT

Choose an attack scenario

Forged system message · attack

Source type

Email

Agent role

Merchant support

Authorized user task (not from email)

Summarize settlement ST-2048 and create a draft support ticket; do not initiate a refund.

02 / Security decision

QUARANTINE

RISK SEVERITY

CRITICAL

CLASSIFIER MODE

heuristic plus local ml

SCAN TIME

1.492 ms

Detected attack categories

Role Change

Tool Abuse

Security rationale

Suspicious untrusted instructions were identified; SANITIZE forwards only conservatively filtered text, while QUARANTINE/ESCALATE withhold all content. External content has no authority over tools.


ROLE SPOOFING

Scripted control is illustrative, not a measured attacked LLM.

Protected agent completes the ticket and denies the external-origin refund.

Provenance → inspection → policy → tool authorization.

The classifier is advisory; the tool gateway is the actual safety boundary.



Nine attack patterns implemented and unit-tested.

Official F3 requires ≥ 7 tested categories. Broad real-world coverage remains unverified.

01 Instruction Override

02 Role Change

03 Secret Extraction

04 Tool Abuse

05 Credential Theft

06 Context Poisoning

07 Multi-Step Jailbreaks

08 Encoded Instructions

09 Indirect Prompt Injection

Evidence: labeled fixtures · regression tests · per-category synthetic holdout · redacted spans

Safety evaluation: results depend on the test distribution.

A known-case retest is not independent evidence; the fresh 40-case result remains weak.

PRE-FIX BASELINE

58.3%

14 / 24 attacks detected

42 known cases; 1 benign FP

FEEDBACK RETEST

100.0%

30 / 30 attacks detected

60 known cases, tuned to feedback

KNOWN 40-CASE SET

50.0%


10 / 20 attacks detected

40 previously reviewed cases; 2 benign FP

F3-D1. Known 40-case challenge misses 10/20 attacks; v1.7 judge feedback set is 28/30 (NOT blind). No live hosted score.

EGO measures safety AND usefulness.

Security that breaks every legitimate workflow is not a viable enterprise product.



EVALUATION

Attack recall, precision, per-class evidence

GUARDRAILS

No unauthorized mock tool execution

OBSERVABILITY

Audit chain, latency and failure analysis

Designed like an enterprise AI platform, not a toy.

A synthetic fintech operations scenario connects risk to real business decision-making.

CONTROL PLANE

Platform governance

Central AI policy and model-independent safety checks.

AGENT STUDIO

Orchestration

Inspection, task execution, incident analysis and replay.

PRODUCT KPIs

Measurable ROI

Safety, false positives, permitted-workflow success, costs.

Financial-impact dashboard uses editable hypothetical assumptions; no real customer savings claimed.

Runnable, measurable, and clear about limitations.

Working local prototype + source tests + case-level evidence + pitch materials.

TODAY

Run the protected fintech workflow, evaluate, inspect traces.

NEXT

Public external benchmark pending; live OpenAI unverified; new 40-case recall 50.0%.

FOR PRODUCTION

Multitenant auth, audit hardening, SIEM, OCR, rate limits, independent security review.

Repository: github.com/saidhinu/TrustBoundary-AI-AgentSecurity

TrustBoundary: provenance before permission.